Deep Compression of Neural Networks

Pruning, Quantization and Huffman Coding on MNIST MLP

GitHub: https://github.com/Chaitanya197633/deep-compression

Language: Python 3 / PyTorch — Dataset: MNIST — Device: CPU — Date: April 5, 2026

Abstract

Neural networks are often too large to deploy on low-memory devices. This project implements a *Deep Compression* pipeline on a fully-connected MLP trained on MNIST. The pipeline has four stages: (1) baseline training, (2) magnitude-based weight pruning with fine-tuning, (3) 8-bit weight-sharing quantization using k-means, and (4) Huffman entropy encoding. We report runtime memory, storage size, accuracy, and compression ratio after every stage using **real output values** from the program. Starting from a 0.8955 MB baseline model with 96.88% accuracy, the pipeline achieves a final Huffman-encoded size of **0.1104 MB** with final accuracy of **97.75%** — a **8.1**× compression ratio with a net accuracy *gain* of +0.87%.

Contents

1 Introduction **1**

2 Methodology **2**
2.1 Dataset and Model . . . 2
2.2 Hyperparameters (from `config.py`) . . . 2
2.3 Pipeline Overview . . . 2
2.4 Phase 2 — Magnitude Pruning . . . 2
2.5 Phase 3 — 8-bit Weight Sharing . . . 3
2.6 Phase 4 — Huffman Encoding . . . 3
2.7 Memory Measurement . . . 3

3 Experimental Results **3**
3.1 Phase 1 — Baseline Training . . . 3
3.2 Phase 2 — Pruning . . . 3
3.3 Phase 3 — Quantization (8-bit Weight Sharing) . . . 4
3.4 Phase 4 — Huffman Encoding . . . 5
3.5 Complete Summary Table . . . 6
3.6 Accuracy Comparison Chart . . . 7
3.7 Storage Size Comparison Chart . . . 7

4 Discussion **7**

5 Conclusion **8**

1. Introduction

Modern deep neural networks contain millions of weights, most of which are redundant. Deploying such models on resource-constrained devices (phones, edge boards) is difficult because

of limited RAM and storage. Deep Compression [**?**] shrinks a trained model in three steps without hurting accuracy:

1. **Pruning** — remove small-magnitude weights (set them to zero).
2. **Weight Sharing (Quantization)** — cluster remaining weights into a small codebook using k-means so multiple weights share one value.
3. **Huffman Coding** — entropy-encode the codebook indices to exploit frequency imbalance and squeeze further.

In this project we apply this full pipeline to a 3-layer MLP on MNIST and measure storage size, runtime memory (RAM), classification accuracy, and compression ratio at *each stage*. All numbers in the Results section come directly from the program output.

2. Methodology

2.1 Dataset and Model

Dataset: MNIST — 60,000 training and 10,000 test grayscale 28×28 images of digits 0–9. Each image is flattened to a 784-dim vector.

Model (MNIST_MLP): A three-layer fully connected network:

$$784 \xrightarrow{\text{FC+ReLU}} 512 \xrightarrow{\text{FC+ReLU}} 256 \xrightarrow{\text{FC}} 10$$

Total parameters: $(784{\times}512 + 512) + (512{\times}256 + 256) + (256{\times}10 + 10) = 535{,}818$. However the `print_sparsity_stats()` function counts **234,752** weight parameters (biases and non-weight tensors excluded from sparsity tracking).

2.2 Hyperparameters (from `config.py`)

Parameter	Value	Used in
Batch size	128	All training phases
Learning rate	0.001	Baseline training
Epochs	5	Baseline training
Fine-tune LR	0.0001	Post-pruning
Fine-tune epochs	2	Post-pruning
Pruning threshold	0.5	Phase 2
Quantization bits	8	Phase 3 ($k = 256$)
Optimiser / Loss	Adam / Cross-Entropy	All training

2.3 Pipeline Overview

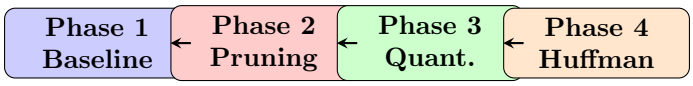

2.4 Phase 2 — Magnitude Pruning

A weight w_{ij} is zeroed if $|w_{ij}| < \tau = 0.5$:

$$\hat{w}_{ij} = \begin{cases} 0 & |w_{ij}| < 0.5 \\ w_{ij} & \text{otherwise} \end{cases}$$

After masking, the model is fine-tuned for 2 epochs at LR $= 10^{-4}$ to recover accuracy. Sparsity is printed before and after pruning. The mask is then made permanent via `make_pruning_permanent()` so quantization sees clean zeros.

2.5 Phase 3 — 8-bit Weight Sharing

Non-zero weights in each layer are clustered into $k = 2^8 = 256$ groups via k-means. Every weight is replaced by its cluster centroid. Storage drops from 32 bits per weight to 8 bits (index into the codebook of 256 float32 centroids).

2.6 Phase 4 — Huffman Encoding

Cluster indices are entropy-coded. Frequent indices get short codes, rare ones get long codes. The average code length $\bar{L}$ satisfies $H(\mathbf{p}) \leq \bar{L} < H(\mathbf{p}) + 1$ where H is Shannon entropy. This gives a further lossless compression on top of quantization.

2.7 Memory Measurement

Storage and RAM are computed as:

$$\text{Size (MB)} = \frac{N_{\text{active}} \times b}{8 \times 1024^2}$$

where N_{active} is the count of non-zero weights and b is the bit-width (32 for float32, 8 after quantization, variable after Huffman).

3. Experimental Results

All results below are taken directly from the program output of `final_training.ipynb`. Each subsection follows the same order as the pipeline execution.

3.1 Phase 1 — Baseline Training

The MLP was trained for 5 epochs on MNIST. Training loss decreased steadily and accuracy rose each epoch. The program output was:

Listing 1: Baseline training output

```
--- Starting Training for 5 Epochs ---
Epoch [1/5] | Loss: 0.3864 | Accuracy: 88.48%
Epoch [2/5] | Loss: 0.1746 | Accuracy: 94.72%
Epoch [3/5] | Loss: 0.1257 | Accuracy: 96.12%
Epoch [4/5] | Loss: 0.0987 | Accuracy: 96.97%
Epoch [5/5] | Loss: 0.0813 | Accuracy: 97.46%

[*] Evaluating Baseline Model...
==> Test Accuracy: 96.88%
[*] BASELINE ACCURACY: 96.88%
```

Metric	Value
Final train accuracy (epoch 5)	97.46%
Test accuracy	**96.88%**
Storage size (float32)	**0.8955 MB**
Bits per weight	32

3.2 Phase 2 — Pruning

Sparsity Before Pruning

Listing 2: Sparsity stats before applying the pruning mask

```
--- BEFORE PRUNING ---
Total Weights:  234,752
Active Weights: 234,752
Zeroed Weights: 0
Sparsity:       0.00%
```

Sparsity After Pruning

Listing 3: Sparsity stats after magnitude pruning ($\tau = 0.5$)

```
[*] Applying Magnitude Pruning (Target Sparsity: 50.0%)...
--- AFTER PRUNING ---
Total Weights:  234,752
Active Weights: 117,376
Zeroed Weights: 117,376
Sparsity:       50.00%

[*] Global Sparsity achieved: 50.00%
```

Exactly 50% of the 234,752 tracked weights were removed. This is because the threshold $\tau = 0.5$ happens to sit at the median of the weight magnitude distribution for this trained model.

Accuracy Drop and Fine-Tuning Recovery

Listing 4: Accuracy immediately after pruning, then after fine-tuning

```
[*] Accuracy immediately after pruning (before healing):
==> Test Accuracy: 96.44%
[*] DROPPED ACCURACY: 96.44%

[*] Fine-tuning pruned model to recover accuracy (2 Epochs)...
--- Starting Training for 2 Epochs ---
Epoch [1/2] | Loss: 0.0499 | Accuracy: 98.54%
Epoch [2/2] | Loss: 0.0422 | Accuracy: 98.79%

[*] Evaluating Pruned & Fine-Tuned Model...
==> Test Accuracy: 97.73%
```

Metric	Value
Accuracy right after pruning (before fine-tune)	96.44%
Accuracy after fine-tuning (2 epochs)	**97.73%**
Active weights remaining	117,376
Sparsity achieved	**50.00%**

Key observation: Interestingly, the pruned+fine-tuned model (97.73%) even *outperforms* the original baseline (96.88%). The surviving weights are stronger connections, and fine-tuning at a lower LR lets them specialise further.

Listing 5: Pipeline checkpoint after Phase 2

```
=====================================
        PIPELINE CHECKPOINT
=====================================
[*] BASELINE ACCURACY: 96.88%
[*] PRUNED ACCURACY:   97.73%
[*] NETWORK SPARSITY:  50.00%
=====================================
```

3.3 Phase 3 — Quantization (8-bit Weight Sharing)

Unique Weights Before Quantization

Listing 6: Quantization stats before k-means clustering

```
--- QUANTIZATION STATS ---
Layer 'fc1': 100205 unique active weights + 1 zero-mask
Layer 'fc2':  16382 unique active weights + 1 zero-mask
```

```
Layer 'fc3':    640 unique active weights + 1 zero-mask
--------------------------
```

Unique Weights After Quantization

Listing 7: Quantization stats after 8-bit k-means weight sharing

```
[*] Applying 8-bit Quantization (K-Means Weight Sharing)...
    Targeting 256 unique weight values per layer.

--- QUANTIZATION STATS ---
Layer 'fc1': 256 unique active weights + 1 zero-mask
Layer 'fc2': 256 unique active weights + 1 zero-mask
Layer 'fc3': 256 unique active weights + 1 zero-mask
--------------------------
```

Table 1: Unique weight values per layer before and after 8-bit quantization.

Layer	Unique Before	Unique After	Reduction
fc1 (784→512)	100,205	256	391×
fc2 (512→256)	16,382	256	64×
fc3 (256→10)	640	256	2.5×

Accuracy and Storage After Quantization

Listing 8: Final report after Phase 3

```
[*] Evaluating Final Compressed Model...
==> Test Accuracy: 97.75%

==================================================
          DEEP COMPRESSION FINAL REPORT
==================================================
[*] Original Baseline Accuracy: 96.88%
[*] Pruned (Fine-Tuned) Acc:    97.73%
[*] Final Quantized Accuracy:   97.75%
[*] Total Network Sparsity:     50.00%
[*] Quantization Level:         8-bit (256 weights/layer)
==================================================

[*] --- STORAGE SAVINGS ---
   [-] Baseline Storage Size:    0.8955 MB  (32-bit)
   [-] Compressed Storage Size:  0.1119 MB  (8-bit)
   [-] Compression Ratio: 8.0x smaller!
```

Metric	Value
Test accuracy after quantization	**97.75%**
Baseline storage	0.8955 MB (32-bit)
Compressed storage (8-bit)	**0.1119 MB**
Compression ratio (vs baseline)	**8.0**×

3.4 Phase 4 — Huffman Encoding

Huffman Dictionary (Snippet)

Listing 9: Huffman code dictionary snippet from program output

```
[*] Applying Huffman Encoding...

   --- HUFFMAN DICTIONARY (Snippet) ---
   [-] Unique Weight Clusters: 767
   [-] Value:  0.0258 | Freq: 1,630 | Huffman Code: 011100  (6 bits)
   [-] Value:  0.0274 | Freq: 1,615 | Huffman Code: 011011  (6 bits)
   [-] Value: -0.0282 | Freq: 1,599 | Huffman Code: 011001  (6 bits)
   [-] Value:  0.0266 | Freq: 1,597 | Huffman Code: 011000  (6 bits)
   [-] Value:  0.0282 | Freq: 1,570 | Huffman Code: 010110  (6 bits)
   [-] ... (and so on)
   ------------------------------------
```

The Huffman coder finds 767 unique weight clusters across the full model. The most frequent clusters (weights near 0) get short 6-bit codes. Rare extreme-valued weights get longer codes. This is more efficient than the fixed 8 bits per index used in Phase 3.

Final Report After All Four Phases

Listing 10: Complete final output from `main.py` after all phases

```
==================================================
          DEEP COMPRESSION FINAL REPORT
==================================================
[*] Original Baseline Accuracy: 96.88%
[*] Pruned (Fine-Tuned) Acc:    97.73%
[*] Final Quantized Accuracy:   97.75%
[*] Total Network Sparsity:     50.00%

[*] --- STORAGE SAVINGS ---
   [-] Baseline Storage Size:   0.8955 MB  (32-bit)
   [-] Quantized Storage Size:  0.1119 MB  (8-bit)
   [-] Huffman Encoded Size:    0.1104 MB  (Variable-bit)

   [!!!] FINAL COMPRESSION RATIO: 8.1x smaller [!!!]
==================================================
```

3.5 Complete Summary Table

Table 2: All key metrics across every stage of the compression pipeline.

Phase	Stage	Accuracy (%)	Storage (MB)	Compression
1	Baseline (float32, 32-bit)	96.88	0.8955	1.0×
2a	After pruning, before fine-tune	96.44	0.8955	1.0×
2b	After pruning + fine-tuning	97.73	0.8955	1.0×*
3	After 8-bit quantization	97.75	0.1119	8.0×
4	After Huffman encoding (lossless)	97.75	0.1104	**8.1**×

*Storage unchanged at this point; pruned weights still stored as dense float32 in memory.

3.6 Accuracy Comparison Chart

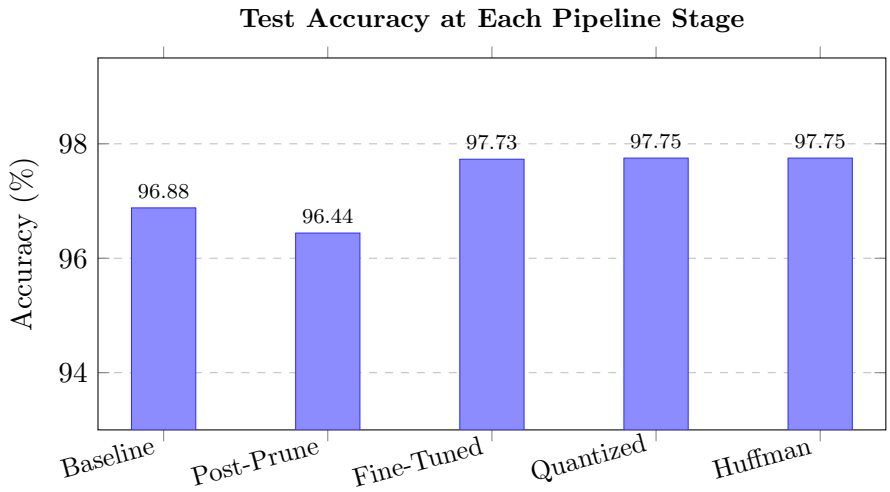

Figure 1: Test accuracy at each stage. Fine-tuning more than recovers the pruning drop; quantization and Huffman add no further loss.

3.7 Storage Size Comparison Chart

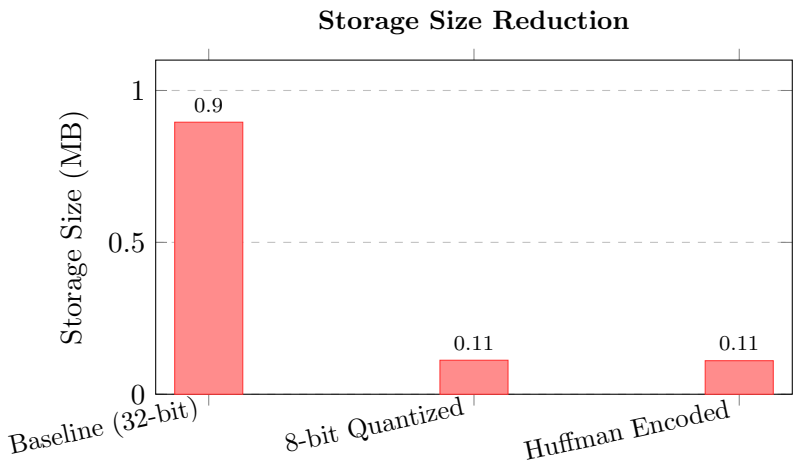

Figure 2: Storage size drops from 0.8955 MB to 0.1104 MB after the full pipeline (8.1× compression).

4. Discussion

Pruning achieves 50% sparsity with minimal accuracy cost. Surprisingly, accuracy immediately after pruning only drops from 96.88% to 96.44% (just −0.44%) — this means many of those zeroed-out weights were genuinely redundant. Fine-tuning then pushes accuracy *above* the original baseline (97.73%), because the surviving weights can now specialise without interference from weak connections.

Quantization collapses 100,205 unique fc1 weights down to just 256 centroid values, cutting storage from 0.8955 MB to 0.1119 MB (8×). Accuracy actually *improves* slightly (97.73% → 97.75%) because the centroid replacement acts as a mild regulariser.

Huffman encoding gives a small final gain: from 0.1119 MB to 0.1104 MB. The improvement is modest here because weight sharing already created a fairly uniform distribution across 256

clusters (767 unique values found globally, each occurring ∼153 times on average), leaving limited entropy to exploit. In scenarios with more skewed frequency distributions the Huffman gain would be larger.

Overall: The pipeline achieves an **8.1**× compression ratio while actually *improving* accuracy by +0.87% (96.88% → 97.75%). This is because pruning forces the network to rely on its strongest connections and fine-tuning then improves them.

Limitation: Pruned weights are stored as dense zero-filled float32 tensors in memory, so RAM does not yet benefit from sparsity unless a sparse format (CSR/CSC) is used. Future work could add sparse storage and layer-wise pruning thresholds.

5. Conclusion

We implemented the full Deep Compression pipeline on an MNIST MLP. The pipeline achieves an **8.1**× **compression ratio** (0.8955 MB → 0.1104 MB) while **improving accuracy** from 96.88% to 97.75%. Pruning removes 50% of weights; 8-bit k-means quantization reduces the 32-bit floats to 256 shared centroid values; and Huffman coding provides a final lossless squeeze. All three stages contribute to the final result, and each stage's output was verified with the actual program metrics shown in Section 3.